

is high value (predicted correct). The second column of the table summarizes the verifier’s prediction, and the third column indicates whether the generated model completion was actually correct or incorrect. Any disagreement between the second and third columns indicates that the verifier made an error.

The first row includes a true positive example, where the verifier correctly classifies the completion as correct. Note that the model is initially unsure about whether the solution is correct and gradually gains certainty as the solution progresses: this is likely a property of the verifier training procedure, where it trains on a large fraction of incorrect model-generated samples.

The second row contains a problem where the solution is correct, but the verifier has rated it as incorrect. This is potentially due to the ambiguity between the “4 times” and the “4 potatoes” in the problem description.

The third row consists of another false negative example. However, unlike the previous example, here the model completion contains some faulty reasoning. As such, even though the final answer in the model completion was correct, the natural language explanation was incorrect, and so the verifier correctly assigned a low score.

In the fourth row we see the verifier score a model completion that starts out correct, but where the verifier gradually becomes less confident in the solution as the solution progresses. After the solution makes a clear mistake (saying that \$64 dollars were spent, instead of the $64 + 16 + 8 = \$88$), the verifier judges the solution as incorrect with a high degree of confidence.

The final row contains a false positive, where the model makes a mistake on the second step, where it subtracts 400 from the price of a diamond jewel instead of a gold one. Verifiers occasionally make mistakes with performing this variable binding of quantities to their relationships.