

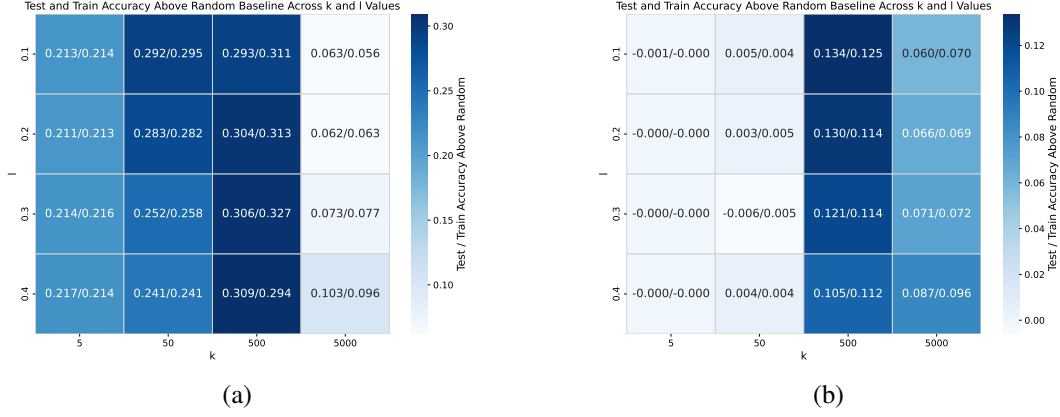


Figure 11: Accuracy gain over a random baseline from linear probes on: (a) Gemma 2 2B, and (b) Pythia 12B. Each cell displays the test/train accuracy above random for a given combination of (k, l) . Darker blue indicates greater accuracy gains.

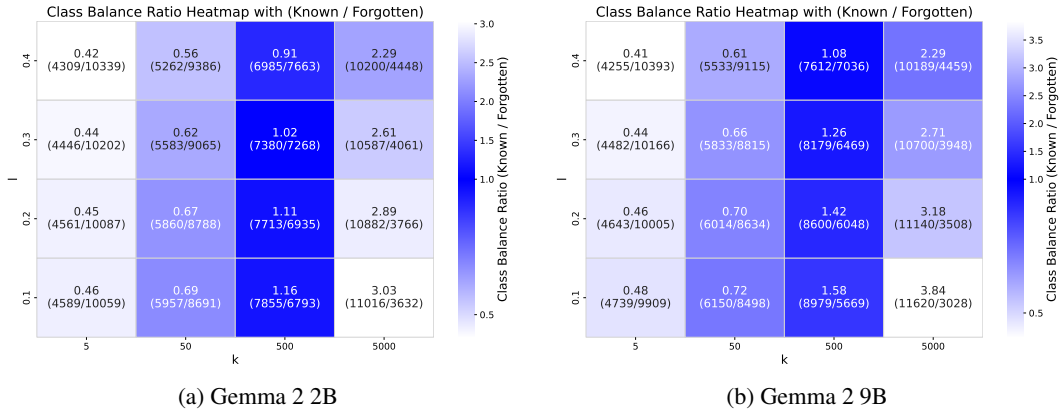
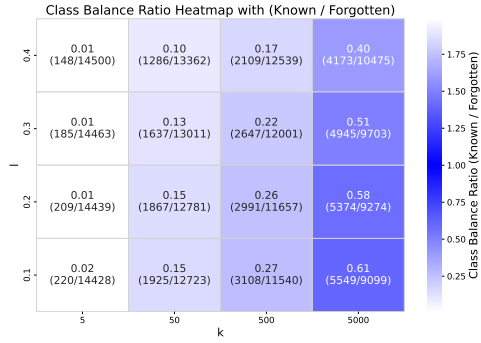
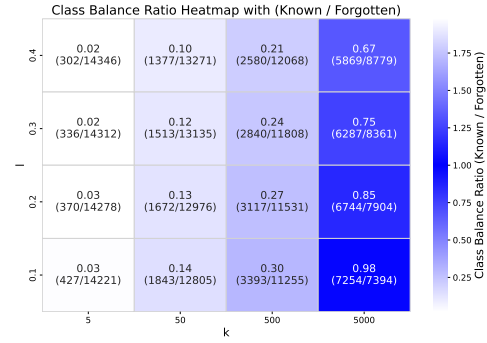


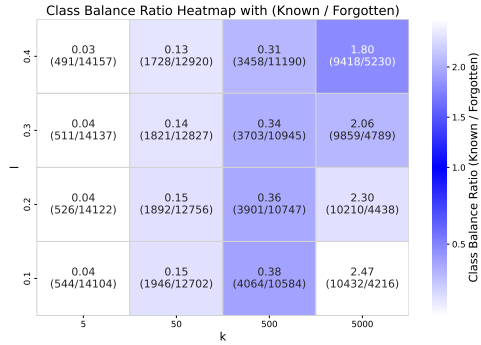
Figure 12: k - l parameters dependence on number of class balance (known and forgotten samples ratio) heatmaps for Gemma 2 models.



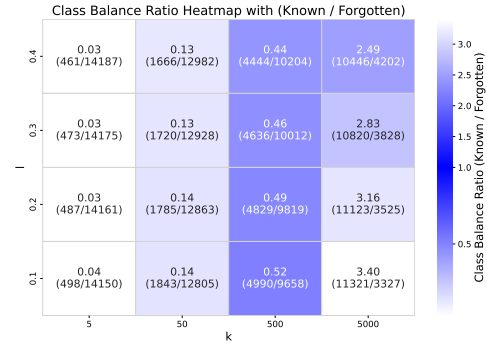
(a) Pythia 70M



(b) Pythia 1.4B



(c) Pythia 6.9B



(d) Pythia 12B

Figure 13: k - l parameters dependence on number of class balance (known and forgotten samples ratio) heatmaps for Pythia models.